



# JSON API Do's & Don'ts

Jan Varljen, Productive CTO



*json:api is a specification for  
building APIs in JSON*

**HTTP://JSONAPI.ORG/**

# CONCEPTS

1. Document structure
2. Fetching
3. CRUD
4. Errors
5. Do's
6. Don'ts

# DOCUMENT STRUCTURE

# JSON STRUCTURE

- data
- included
- links
- meta

## EXAMPLE

<http://api.productive.io.dev/api/v2/1/projects?token=abc>

# RESOURCE

- id
- type
- attributes
- relationships



## EXAMPLE

<http://api.productive.io.dev/api/v2/1/projects/363?token=abc>

# RELATIONSHIPS

- data
- link

# EXAMPLE

```
// ...
{
  "type": "articles",
  "id": "1",
  "attributes": {
    "title": "Rails is Omakase"
  },
  "relationships": {
    "author": {
      "links": {
        "self": "http://example.com/articles/1/relationships/author",
        "related": "http://example.com/articles/1/author"
      },
      "data": { "type": "people", "id": "9" }
    }
  },
  "links": {
    "self": "http://example.com/articles/1"
  }
}
// ...
```

# INCLUDED

- related resources included inside the same response

## EXAMPLE

<http://api.productive.io.dev/api/v2/1/projects/363?token=abc>

# LINKS & META

- pagination links
- pagination metadata
- any other provisional metadata

## EXAMPLE

<http://api.productive.io.dev/api/v2/1/projects?token=abc>

**FETCHING**



## INCLUDING RELATED RESOURCES

- API may return related resources by default
- API may support the `include` request parameter

## INCLUDE PARAMETER

- GET /projects/1?include=company
- GET /projects/1?include=project\_manager
- GET /projects/1?include=company,project\_manager
- GET /projects/1?include=project\_manager.company

## **FILTERING RESOURCES**

- use the `filter` query parameter for filtering resources
- JSON API is agnostic about the strategies supported by a server

## **FILTER PARAMETER**

- GET /projects?filter[company\_id]=41
- [http://api.productive.io.dev/api/v2/1/projects?  
filter%5Bcompany\\_id%5D=41&token=abc](http://api.productive.io.dev/api/v2/1/projects?filter%5Bcompany_id%5D=41&token=abc)

# FRACTIONAL ATTRIBUTES

- API may return only specified attributes in the response
- use the `fields` query parameter for filtering attributes

## FIELDS PARAMETER

- GET /projects/1?fields[projects]=name
- GET /projects/1?fields[projects]=name,currency
- GET /projects/1?fields[projects]=name,currency&fields[company]=name

## **SORTING RESOURCES**

- API may support sorting resources
- use the `sort` query parameter for sorting

## **SORT PARAMETER**

- GET /projects?sort=name
- GET /projects?sort=profit,name
- GET /projects?sort=-profit,name



# PAGINATION

- API endpoint may be paginated
- pagination links are in the `links` part of the response
- <http://api.productive.io.dev/api/v2/1/projects?token=abc>

C(R)UD

# CREATE

- POST /projects
- JSON:API request body
- 201 Created/202 Accepted/204 No content

# EXAMPLE

```
{
  "data": {
    "type": "projects",
    "attributes": {
      "name": "project name"
    },
    "relationships": {
      "company": {
        "data": {
          "type": "companies",
          "id": "1"
        }
      },
      "project_manager": {
        "data": {
          "type": "people",
          "id": "1"
        }
      }
    }
  }
}
```

# UPDATE

- PATCH /projects/1
- JSON:API request body
- partial updates
- 200 OK/202 Accepted/204 No content

# EXAMPLE

```
{  
  "data": {  
    "type": "projects",  
    "id": "1",  
    "attributes": {  
      "name": "new name"  
    },  
    "relationships": {  
      "company": {  
        "data": {  
          "type": "companies",  
          "id": "2"  
        }  
      }  
    }  
  }  
}
```

# DELETE

- DELETE /projects/1
- no request body
- 200 OK/202 Accepted/204 No content

**ERRORS**



# ERRORS STRUCTURE

- errors (array)

# ERROR STRUCTURE

- status
- title
- details
- source
- meta

## EXAMPLE

<http://api.productive.io.dev/api/v2/1/projects/363>

<http://api.productive.io.dev/api/v2/1/projects/99999?token=abc>

# EXAMPLE

```
{
  "errors": [
    {
      "status": "422",
      "title": "Invalid Attribute",
      "detail": "can't be blank",
      "source": {
        "pointer": "data/attributes/name"
      }
    },
    {
      "status": "422",
      "title": "Invalid Attribute",
      "detail": "can't be blank",
      "source": {
        "pointer": "data/attributes/project_manager"
      }
    },
    {
      "status": "422",
      "title": "Invalid Attribute",
      "detail": "can't be blank",
      "source": {
        "pointer": "data/attributes/company"
      }
    }
  ]
}
```



DO'S

# FLAT STRUCTURE

- `/projects/363/comments` vs `/comments?filter[project_id]=363`
- every resource on own endpoint:
  - `people/1/comments`
  - `companies/1/comments`
  - `invoices/1/comments`
  - ....
  - becomes `/comments?filter[x]`

# FLAT STRUCTURE

- less code duplications
- you need /comments route either way:
  - PATCH /comments/1
  - DELETE /comments/1
- more flexible:
  - /comments?filter[project\_id][]=363&filter[project\_id][]=364

# PAGINATE EVERY ENDPOINT

- /projects
  - returns 10 projects
- /projects
  - returns 1000 projects
  - slow, blocking
- always paginate every endpoint



# MAX PER PAGE

- `/projects?page[size]=1_000_000`
  - boom!
- always have a `max_per_page` defined
  - `max_per_page = 100`
- add info to meta
- have a system to override it by resource
  - `/tags` → `max_per_page=300`
  - `/projects` → `max_per_page=50`

# POST CREATE RESPONSE

- return 201 instead of 204 on POST create
- 201 Created
  - resource included in the response
- 204 No content
  - response empty
  - consider the resource sent in the request is accepted by the server
  - almost never identical
  - think about timestamp columns like (created\_at)
- always return 201 if possible

# CUSTOM ACTIONS

- archive project
  - `project.archived_at = 2017-04-18 10:25:49`
- this should be set by the server not client
- PATCH with `archived_at: {time_on_client}` is wrong
- PATCH `/projects/:id/archive`
  - server sets timestamps and returns 200 OK with updated resource

# EXTRACT RESOURCES

- update user password
  - password\_reset\_token
  - password\_token\_sent\_at
  - password\_token\_valid\_for
  - ...
- usually columns in the users table
- PATCH /users/1
  - bunch of “magic” attributes pollute the user resource
  - used only for password reset (rarely)
- API resources != database structure

# EXTRACT RESOURCES

- password becomes first-class citizen → independent resource
- CREATE /passwords
  - create new reset password token
- GET /passwords/:token
  - get a password reset token
- PATCH /passwords/:token
  - make the password change (if token still valid)

# CACHING

- cache every resource individually
- cache key example:
  - `/ {resource_name} / {object_key} / {serializer_file_digest}`
- `/projects/363-201704131447176380000000/7c4eba8cd5a6cd89d6e07ce1a51cc23f`
- `http://api.productive.io.dev/api/v2/1/projects/363?token=abc`

# CACHE READ MULTI

- <http://api.productive.io.dev/api/v2/1/projects?token=abc>
- 30 projects + 17 included resources → 47 cache store hits
- JSON API is agnostic about caching implementation
- write you're own caching layer
  - use read multi from cache (ideally only 1 cache request)
  - solve application specific problems (Productive has roles)

# NANOSECONDS FOR CACHE

- many frameworks use timestamp in seconds for cache
  - “2017-04-18 10:25:26”
- multiple updates in single second – happens all the time
  - invalid responses
- use milliseconds or nanoseconds for cache



# SEPARATION OF CONCERNS

- ProjectsController
  - handles authentication and authorization
- ProjectFilter
  - handles `filter` and `sort` param and filtering
- ProjectSerializer
  - handles `fields` param and fractional attributes
- ProjectPreloader
  - handles `includes` param and preloading relationships
- ProjectForm
  - handles validations and saving resources
- ...



**DON'TS**

# NULLS ON PARTIAL UPDATES

- PATCH request can update a single attribute
- missing attributes and relations should not be treated as NULL

```
{  
  "data": {  
    "type": "projects",  
    "id": "1",  
    "attributes": {  
      "name": "new name"  
    }  
  }  
}
```

# ABUSE DELETE VERB

- archive project
  - DELETE /projects/363
- archive touches a timestamp – does not actually REMOVE the resource
- DELETE is intended to REMOVE the resource
  - response is 204 No content
- archiving using DELETE will work but problems happen afterwards
  - compliant clients will remove the resource

# BREAK CACHE

- resources are cached individually
- misplaced attributes can break cache
  - company\_name

```
▼ {  
  ▼ "data": {  
    "id": "363",  
    "type": "projects",  
    ▼ "attributes": {  
      "name": "AAA",  
      "company_name": "Mariplast"  
    },  
    ▼ "relationships": {  
      ▼ "company": {  
        ▼ "data": {  
          "type": "companies",  
          "id": "1"  
        }  
      },  
      ▼ "project_manager": {  
        ▼ "data": {  
          "type": "people",  
          "id": "1"  
        }  
      },  
      ▼ "last_actor": {  
        ▼ "data": {  
          "type": "people",  
          "id": "1"  
        }  
      }  
    }  
  },  
  ► "included": [ ... ] // 2 items  
}
```

# BREAK CACHE

- sometimes we need “counters”
- can also break cache if handled improperly
- two important things
  - cache on the database layer
  - manual cache refresh when needed

```
▼ {  
  ▼ "data": {  
    "id": "363",  
    "type": "projects",  
    ▼ "attributes": {  
      "name": "AAA",  
      "open_tasks_count": 28  
    },  
    ► "relationships": { ... } // 3 items  
  },  
  ► "included": [ ... ] // 2 items  
}
```

# NEVER INCLUDE HAS-MANY RELATIONS

- big responses
- pagination problems
- filtering problems
- sorting problems
- prefer separate calls
  - `/tasks?filter[project_id]=11`

```
▼ {  
  ▼ "data": {  
    "id": "11",  
    "type": "projects",  
    ► "attributes": { ... }, // 1 item  
    ▼ "relationships": {  
      ▼ "tasks": {  
        ► "data": [ ... ] // 72 items  
      }  
    }  
  },  
  ► "included": [ ... ] // 72 items  
}
```



# Thank you!

JAN@PRODUCTIVE.IO

Visit productive.io.